

Satisfiability modulo theories

In [computer science](#) and [mathematical logic](#), the **satisfiability modulo theories (SMT)** problem is a [decision problem](#) for logical formulas with respect to combinations of background theories expressed in classical [first-order logic](#) with equality. Examples of theories typically used in computer science are the theory of [real numbers](#), the theory of [integers](#), and the theories of various [data structures](#) such as [lists](#), [arrays](#), [bit vectors](#) and so on. SMT can be thought of as a form of the [constraint satisfaction problem](#) and thus a certain formalized approach to [constraint programming](#).

Contents

- Basic terminology**
- Expressive power of SMT**
- SMT solver approaches**
- SMT for undecidable theories**
- SMT solvers**
- Applications**
 - Verification
 - Symbolic-execution based analysis and testing
- See also**
- Notes**
- References**

Basic terminology

Formally speaking, an SMT instance is a [formula](#) in [first-order logic](#), where some function and predicate symbols have additional interpretations, and SMT is the problem of determining whether such a formula is satisfiable. In other words, imagine an instance of the [Boolean satisfiability problem](#) (SAT) in which some of the binary variables are replaced by [predicates](#) over a suitable set of non-binary variables. A predicate is basically a binary-valued function of non-binary variables. Example predicates include linear [inequalities](#) (e.g., $3x + 2y - z \geq 4$) or equalities involving [uninterpreted terms](#) and function symbols (e.g., $f(f(u, v), v) = f(u, v)$ where f is some unspecified function of two arguments.) These predicates are classified according to each respective theory assigned. For instance, linear inequalities over real variables are evaluated using the rules of the theory of linear real [arithmetic](#), whereas predicates involving uninterpreted terms and function symbols are evaluated using the rules of the theory of [uninterpreted functions](#) with equality (sometimes referred to as the [empty theory](#)). Other theories include the theories of [arrays](#) and [list structures](#) (useful for modeling and verifying [computer programs](#)), and the theory of [bit vectors](#) (useful in modeling and verifying [hardware designs](#)). Subtheories are also possible: for example, difference logic is a sub-theory of linear arithmetic in which each inequality is restricted to have the form $x - y > c$ for variables x and y and constant c .

Most SMT solvers support only [quantifier free](#) fragments of their logics.

Expressive power of SMT

An SMT instance is a generalization of a Boolean SAT instance in which various sets of variables are replaced by [predicates](#) from a variety of underlying theories. Obviously, SMT formulas provide a much richer [modeling language](#) than is possible with Boolean SAT formulas. For example, an SMT formula allows us to model the [datapath](#) operations of a [microprocessor](#) at the word rather than the bit level.

By comparison, [answer set programming](#) is also based on predicates (more precisely, on [atomic sentences](#) created from [atomic formula](#)). Unlike SMT, answer-set programs do not have quantifiers, and cannot easily express constraints such as [linear arithmetic](#) or [difference logic](#)—ASP is at best suitable for boolean problems that reduce to the [free theory](#) of uninterpreted functions. Implementing 32-bit integers as bitvectors in ASP suffers from most of the same problems that early SMT solvers faced: "obvious" identities such as $x+y=y+x$ are difficult to deduce.

[Constraint logic programming](#) does provide support for linear arithmetic constraints, but within a completely different theoretical framework.

SMT solver approaches

Early attempts for solving SMT instances involved translating them to Boolean SAT instances (e.g., a 32-bit integer variable would be encoded by 32 bit variables with appropriate weights and word-level operations such as 'plus' would be replaced by lower-level logic operations on the bits) and passing this formula to a Boolean SAT solver. This approach, which is referred to as *the eager approach*, has its merits: by pre-processing the SMT formula into an equivalent Boolean

SAT formula we can use existing Boolean SAT solvers "as-is" and leverage their performance and capacity improvements over time. On the other hand, the loss of the high-level semantics of the underlying theories means that the Boolean SAT solver has to work a lot harder than necessary to discover "obvious" facts (such as $x + y = y + x$ for integer addition.) This observation led to the development of a number of SMT solvers that tightly integrate the Boolean reasoning of a DPLL-style search with theory-specific solvers (*T-solvers*) that handle conjunctions (ANDs) of predicates from a given theory. This approach is referred to as *the lazy approach*.

Dubbed DPLL(T),^[1] this architecture gives the responsibility of Boolean reasoning to the DPLL-based SAT solver which, in turn, interacts with a solver for theory T through a well-defined interface. The theory solver only needs to worry about checking the feasibility of conjunctions of theory predicates passed on to it from the SAT solver as it explores the Boolean search space of the formula. For this integration to work well, however, the theory solver must be able to participate in propagation and conflict analysis, i.e., it must be able to infer new facts from already established facts, as well as to supply succinct explanations of infeasibility when theory conflicts arise. In other words, the theory solver must be incremental and backtrackable.

SMT for undecidable theories

Most of the common SMT approaches support decidable theories. However, many real-world systems can only be modelled by means of non-linear arithmetic over the real numbers involving transcendental functions, e.g. an aircraft and its behavior. This fact motivates an extension of the SMT problem to non-linear theories, e.g. determine whether

$$(\sin(x))^3 = \cos(\log(y) \cdot x) \vee b \vee -x^2 \geq 2.3y) \wedge (\neg b \vee y < -34.4 \vee \exp(x) > \frac{y}{x})$$

where

$$b \in \mathbb{B}, x, y \in \mathbb{R}$$

is satisfiable. Then, such problems become undecidable in general. (It is important to note, however, that the theory of real closed fields, and thus the full first order theory of the real numbers, are decidable using quantifier elimination. This is due to Alfred Tarski.) The first order theory of the natural numbers with addition (but not multiplication), called Presburger arithmetic, is also decidable. Since multiplication by constants can be implemented as nested additions, the arithmetic in many computer programs can be expressed using Presburger arithmetic, resulting in decidable formulas.

Examples of SMT solvers addressing Boolean combinations of theory atoms from undecidable arithmetic theories over the reals are ABSolver,^[2] which employs a classical DPLL(T) architecture with a non-linear optimization packet as (necessarily incomplete) subordinate theory solver, and iSAT ^[1] (<http://isat.gforge.avac.s.org/>), building on a unification of DPLL SAT-solving and interval constraint propagation called the iSAT algorithm.^[3]

SMT solvers

The table below summarizes some of the features of the many available SMT solvers. The column "SMT-LIB" indicates compatibility with the SMT-LIB language; many systems marked 'yes' may support only older versions of SMT-LIB, or offer only partial support for the language. The column "CVC" indicates support for the CVC language. The column "DIMACS" indicates support for the DIMACS (<http://www.satcompetition.org/2009/format-benchmarks2009.html>) format.

Projects differ not only in features and performance, but also in the viability of the surrounding community, its ongoing interest in a project, and its ability to contribute documentation, fixes, tests and enhancements.

Platform			Features						Notes
Name	OS	License	SMT-LIB	CVC	DIMACS	Built-in theories	API	SMT-COMP [2] (http://www.smtcomp.org/)	
ABsolver (http://absolver.sourceforge.net/)	Linux	CPL	v1.2	No	Yes	linear arithmetic, non-linear arithmetic	C++	no	DPLL-based
Alt-Ergo	Linux , Mac OS , Windows	CeCILL-C (roughly equivalent to LGPL)	partial v1.2 and v2.0	No	No	empty theory , linear integer and rational arithmetic , non-linear arithmetic , polymorphic arrays , enumerated datatypes , AC symbols , bitvectors , record datatypes , quantifiers	OCaml	2008	Polymorphic first-order input language à la ML, SAT-solver based, combines Shostak-like and Nelson-Oppen like approaches for reasoning modulo theories
Barcelogic (http://www.lsi.upc.edu/~oliveras/bcrlt-main.html)	Linux	Proprietary	v1.2			empty theory , difference logic	C++	2009	DPLL-based, congruence closure
Beaver (http://ucld.eecs.berkeley.edu/jha/beaver-dist/beaver.html)	Linux , Windows	BSD	v1.2	No	No	bitvectors	OCaml	2009	SAT-solver based
Boolector (http://fmv.jku.at/boolector/index.html)	Linux	GPLv3	v1.2	No	No	bitvectors , arrays	C	2009	SAT-solver based
CVC3 (http://www.cs.nyu.edu/acsys/cvc3/)	Linux	BSD	v1.2	Yes		empty theory , linear arithmetic , arrays , tuples , types , records , bitvectors , quantifiers	C/C++	2010	proof output to HOL
CVC4 (http://c4.cs.nyu.edu/)	Linux , Mac OS , Windows	BSD	Yes	Yes		rational and integer linear arithmetic , arrays , tuples , records , inductive data types , bitvectors , strings , and equality over uninterpreted function symbols	C++	2010	version 1.4 released July 2014
Decision Procedure Toolkit (DPT) (http://sourceforge.net/projects/dpt)	Linux	Apache	No				OCaml	no	DPLL-based
iSAT (http://isat.gforge.avacs.org/)	Linux	Proprietary	No			non-linear arithmetic		no	DPLL-based
MathSAT (http://mathsat.fbk.eu/)	Linux , Mac OS , Windows	Proprietary	Yes		Yes	empty theory , linear arithmetic , bitvectors , arrays	C/C++ , Python , Java	2010	DPLL-based
MiniSmt (http://cl-informatik.uibk.ac.at/software/minismt/)	Linux	LGPL	partial v2.0			non-linear arithmetic		2010	SAT-solver based, Yices-based
Nom (https://link.springer.com/chapter/10.1007/978-3-319-21690-4_29)									SMT solver for string constraints
OpenCog	Linux	AGPL	No	No	No	probabilistic logic , arithmetic , relational models	C++ , Scheme , Python	no	subgraph isomorphism
OpenSMT (http://)	Linux ,	GPLv3	partial		Yes	empty theory , differences ,	C++	2011	lazy SMT Solver

p://verify.inf.usi.ch/opensmt/	Mac OS, Windows		v2.0			linear arithmetic, bitvectors			
raSAT (http://www.jaist.ac.jp/~s1310007/raSAT/)	Linux	GPLv3	v2.0			real and integer nonlinear arithmetics		2014, 2015	extension of the Interval Constraint Propagation with Testing and the Intermediate Value Theorem
SatEEn (http://vlsi.colorado.edu/~hhkim/sateen/)	?	Proprietary	v1.2			linear arithmetic, difference logic	none	2009	
SMTInterpol (http://ultimate.informatik.uni-freiburg.de/smtinterpol/)	Linux, Mac OS, Windows	LGPLv3	v2.0			uninterpreted functions, linear real arithmetic, and linear integer arithmetic	Java	2012	Focuses on generating high quality, compact interpolants.
SMCHR (http://www.comp.nus.edu.sg/~gregory/smchr/)	Linux, Mac OS, Windows	GPLv3	No	No	No	linear arithmetic, nonlinear arithmetic, heaps	C	no	Can implement new theories using Constraint Handling Rules.
SMT-RAT (https://github.com/smt-rat/smt-rat/wiki)	Linux, Mac OS	MIT	v2.0	No	No	linear arithmetic, nonlinear arithmetic	C++	2015	Toolbox for strategic and parallel SMT solving consisting of a collection of SMT compliant implementations.
SONOLAR (http://www.informatik.uni-bremen.de/agbs/florian/sonolar/)	Linux, Windows	Proprietary	partial v2.0			bitvectors	C	2010	SAT-solver based
Spear (http://www.cs.ubc.ca/~babic/index_spear.htm)	Linux, Mac OS, Windows	Proprietary	v1.2			bitvectors		2008	
STP (http://stp.github.io/)	Linux, OpenBSD, Windows, Mac OS	MIT	partial v2.0	Yes	No	bitvectors, arrays	C, C++, Python, OCaml, Java	2011	SAT-solver based
SWORD (http://www.informatik.uni-bremen.de/agra/eng/sword.php)	Linux	Proprietary	v1.2			bitvectors		2009	
UCLID (http://uclid.eecs.berkeley.edu/wiki/index.php/Main_Page)	Linux	BSD	No	No	No	empty theory, linear arithmetic, bitvectors, and constrained lambda (arrays, memories, cache, etc.)		no	SAT-solver based, written in Moscow ML. Input language is SMV model checker. Well-documented!
veriT (http://www.verit-solver.org/)	Linux, OS X	BSD	partial v2.0			empty theory, rational and integer linear arithmetics, quantifiers, and equality over uninterpreted function symbols	C/C++	2010	SAT-solver based
Yices (http://yices.csl.sri.com/)	Linux, Mac OS, Windows	GPLv3	v2.0	No	Yes	rational and integer linear arithmetic, bitvectors, arrays, and equality over uninterpreted function symbols	C	2014	Source code is available online
Z3 (https://github.com/Z3Prover/z3)	Linux, Mac OS, Windows, FreeBSD	MIT	v2.0		Yes	empty theory, linear arithmetic, nonlinear arithmetic, bitvectors, arrays, datatypes, quantifiers, strings	C/C++, .NET, OCaml, Python, Java, Haskell	2011	

Applications

SMT solvers are useful both for verification, proving the correctness of programs, software testing based on symbolic execution, and for synthesis, generating program fragments by searching over the space of possible programs.

Verification

Computer-aided verification of computer programs often uses SMT solvers. A common technique is to translate preconditions, postconditions, loop conditions, and assertions into SMT formulas in order to determine if all properties can hold.

There are many verifiers built on top of the Z3 (<https://github.com/Z3Prover/z3>) SMT solver. Boogie (<http://research.microsoft.com/en-us/projects/boogie/>) is an intermediate verification language that uses Z3 to automatically check simple imperative programs. The [3] (<http://vcc.codeplex.com>) verifier for concurrent C uses Boogie, as well as Dafny (<http://research.microsoft.com/en-us/projects/dafny/>) for imperative object-based programs, Chalice (<http://research.microsoft.com/en-us/projects/chalice/>) for concurrent programs, and Spec# (<http://research.microsoft.com/en-us/projects/specsharp/>) for C#. F* (<http://research.microsoft.com/en-us/projects/fstar/>) is a dependently typed language that uses Z3 to find proofs; the compiler carries these proofs through to produce proof-carrying bytecode. The sbv (<https://hackage.haskell.org/package/sbv>) library provides SMT-based verification of Haskell programs, and lets the user choose among a number of solvers such as Z3, ABC, Boolector, CVC4, MathSAT and Yices.

There are also many verifiers built on top of the Alt-Ergo (<http://alt-ergo.ocamlpro.com/>) SMT solver. Here is a list of mature applications:

- Why3 (<http://why3.lri.fr/>), a platform for deductive program verification, uses Alt-Ergo as its main prover;
- CAVEAT, a C-verifier developed by CEA and used by Airbus; Alt-Ergo was included in the qualification DO-178C of one of its recent aircraft;
- Frama-C, a framework to analyse C-code, uses Alt-Ergo in the Jessie and WP plugins (dedicated to "deductive program verification");
- SPARK, uses CVC4 and Alt-Ergo (behind GNATprove) to automate the verification of some assertions in SPARK 2014;
- Atelier-B can use Alt-Ergo instead of its main prover (increasing success from 84% to 98% on the ANR Bware project benchmarks (<http://alt-ergo.lri.fr/documents/ABZ-2014.pdf>));
- Rodin, a B-method framework developed by Systerel, can use Alt-Ergo as a back-end;
- Cubicle (<http://cubicle.lri.fr/>), an open source model checker for verifying safety properties of array-based transition systems.
- EasyCrypt (<https://www.easycrypt.info/>), a toolset for reasoning about relational properties of probabilistic computations with adversarial code.

Many SMT solvers implement a common interface format called SMTLIB2 (<http://smt-lib.org/>) (such files usually have the extension ".smt2"). The LiquidHaskell (<https://ucsd-progsys.github.io/liquidhaskell-blog/>) tool implements a refinement type based verifier for Haskell that can use any SMTLIB2 compliant solver, e.g. CVC4, MathSat, or Z3.

Symbolic-execution based analysis and testing

An important application of SMT solvers is symbolic execution for analysis and testing of programs (e.g., concolic testing), aimed particularly at finding security vulnerabilities. Important actively-maintained tools in this category include SAGE (http://research.microsoft.com/en-us/um/people/pg/public_psfiles/ndss2008.pdf) from Microsoft Research, KLEE (<https://klee.github.io/>), S2E (<http://s2e.epfl.ch/>), and Triton (<https://triton.quarkslab.com>). SMT solvers that are particularly useful for symbolic-execution applications include Z3 (<https://github.com/Z3Prover/z3>), STP (<https://sites.google.com/site/stpfastprover/>), Z3str2 (<https://sites.google.com/site/z3strsolver/>), and Boolector (<http://fmv.jku.at/boolector/>).

See also

- Answer set programming

Notes

1. Nieuwenhuis, R.; Oliveras, A.; Tinelli, C. (2006), "Solving SAT and SAT Modulo Theories: From an Abstract Davis-Putnam-Logemann-Loveland Procedure to DPLL(T)", *Journal of the ACM* (<http://homepage.cs.uiowa.edu/~tinelli/papers/NieOT-JACM-06.pdf>) (PDF), **53** (6), pp. 937–977.
2. Bauer, A.; Pister, M.; Tautschnig, M. (2007), "Tool-support for the analysis of hybrid systems and models", *Proceedings of the 2007 Conference on Design, Automation and Test in Europe (DATE'07)*, IEEE Computer Society, p. 1, doi:10.1109/DATE.2007.364411 (<https://doi.org/10.1109%2FDATE.2007.364411>)
3. Fränzle, M.; Herde, C.; Ratschan, S.; Schubert, T.; Teige, T. (2007), "Efficient Solving of Large Non-linear Arithmetic Constraint Systems with Complex Boolean Structure", *JSAT Special Issue on SAT/CP Integration* (http://jsat.ewi.tudelft.nl/content/volume1/JSAT1_11_Fraenzle.pdf) (PDF), **1**, pp. 209–236

References

- C Barrett, R Sebastiani, S Seshia, and C Tinelli, "Satisfiability Modulo Theories." In Handbook of Satisfiability, vol. 185 of Frontiers in Artificial Intelligence and Applications, (A Biere, M J H Heule, H van Maaren, and T Walsh, eds.), IOS Press, Feb. 2009, pp. 825–885.
- Vijay Ganesh (PhD. Thesis 2007), Decision Procedures for Bit-Vectors, Arrays and Integers (https://ece.uwaterloo.ca/~vganesh/Publications_files/vg2007-PhD-STANFORD.pdf), Computer Science Department, Stanford University, Stanford, CA, U.S., Sept 2007
- Susmit Jha, Rishikesh Limaye, and Sanjit A. Seshia. Beaver: Engineering an efficient SMT solver for bit-vector arithmetic. (https://dx.doi.org/10.1007/978-3-642-02658-4_53) In Proceedings of 21st International Conference on Computer-Aided Verification, pp. 668–674, 2009.
- R. E. Bryant, S. M. German, and M. N. Velev, "Microprocessor Verification Using Efficient Decision Procedures for a Logic of Equality with Uninterpreted Functions," in Analytic Tableaux and Related Methods, pp. 1–13, 1999.
- M. Davis and H. Putnam, A Computing Procedure for Quantification Theory (<https://dx.doi.org/10.1145%2F321033.321034>), Journal of the Association for Computing Machinery, vol. 7, no., pp. 201–215, 1960.

- M. Davis, G. Logemann, and D. Loveland, *A Machine Program for Theorem-Proving* (<https://dx.doi.org/10.1145%2F368273.368557>), Communications of the ACM, vol. 5, no. 7, pp. 394–397, 1962.
- D. Kroening and O. Strichman, *Decision Procedures – an algorithmic point of view* (2008), Springer (Theoretical Computer Science series) ISBN 978-3-540-74104-6.
- G.-J. Nam, K. A. Sakallah, and R. Rutenbar, *A New FPGA Detailed Routing Approach via Search-Based Boolean Satisfiability* (<https://dx.doi.org/10.1109%2FTCAD.2002.1004311>), IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems, vol. 21, no. 6, pp. 674–684, 2002.
- SMT-LIB: The Satisfiability Modulo Theories Library (<http://smtlib.org/>)
- SMT-COMP: The Satisfiability Modulo Theories Competition (<http://www.smtcomp.org>)
- Decision procedures - an algorithmic point of view (<http://www.decision-procedures.org/>)
- Summer School on SAT/SMT solvers and their applications (<http://people.csail.mit.edu/vganesh/summerschool>)
- R. Sebastiani, *Lazy Satisfiability Modulo Theories* (http://jsat.ewi.tudelft.nl/content/volume3/JSAT3_9_Sebastiani.pdf), Dipartimento di Ingegneria e Scienza dell'Informazione, Universita di Trento, Italy, Dec 2007
- D.Yurichev, *Quick introduction into SAT/SMT solvers and symbolic execution* (https://yurichev.com/writings/SAT_SMT_draft-EN.pdf)

This article is adapted from a column in the ACM SIGDA (<http://www.sigda.org>) e-newsletter (<http://www.sigda.org/newsletter/index.html>) by Prof. Karem Sakallah (<http://www.eecs.umich.edu/~karem>). Original text is available here (<http://archive.sigda.org/newsletter/2006/061215.txt>)

Retrieved from "https://en.wikipedia.org/w/index.php?title=Satisfiability_modulo_theories&oldid=807092496"

This page was last edited on 25 October 2017, at 21:14.

Text is available under the [Creative Commons Attribution-ShareAlike License](#); additional terms may apply. By using this site, you agree to the [Terms of Use](#) and [Privacy Policy](#). Wikipedia® is a registered trademark of the [Wikimedia Foundation, Inc.](#), a non-profit organization.